

# Data Intensive processing with iRODS and the middleware CiGri for the Whisper project

Briand X.\*

Bzeznik B.†

## Abstract

*blabla*

*blabla*

## Keywords

Data-Intensive, grid computing, distributed storage, Seismic Noise, Whisper, Cigri, Irods.

## 1 Plan

1. Introduction, abstract (3§, 1 page)
2. Background (7-10§, 1-2 pages)
  - (a) The Whisper project
  - (b) Size of observational data
  - (c) Size of computational data
  - (d) General Big data/ science of universe
  - (e) Coupling scientific objectives / IT constraints
  - (f) Data grid paradigm
  - (g) Data grid in Grenoble
  - (h) Collaboration Whisper/IT infrastructure
3. Software for data-intensive processing (6-12§, 1-2 pages)
  - (a) General IT codes whisper
  - (b) Computer language and librairies
  - (c) Structure of codes
  - (d) Package for raw data
  - (e) Package for correlations
  - (f) First optimisation of correlation
  - (g) Second optimisation of correlation
  - (h) Codes for analysis of correlations
  - (i) Where we can run the codes
4. IT infrastrcuture for grid computing (6-12§, 1-2 pages)
  - (a) Needs of coupling storage computation
  - (b) Presentation Ciment
  - (c) Coupling Irods/Cigri
  - (d) Resulting Data-Grid
  - (e) Presentation Irods
  - (f) General Infra Irods (diagram)
  - (g) Effective Nodes Irods (diagram)
  - (h) General Cigri
  - (i) Mechanism with OAR
  - (j) Mechanism resubmission/besteffort
  - (k) Description of a campaign
  - (l) Low Interaction Cigri/Irods
  - (m) Cigri V2 and V3 functionalities
5. Results (28-31§, 6-7 pages)
6. Discussion (7-10§, 1-2 pages)

---

\*Isterre, Cnrs, email xav.briand@gmail.com

†Ciment, Université Joseph Fourier

## 2 Background 7-10§, 1-2 pages

The *Whisper* \* project is a European project on seismology whose goal is to study properties of the earth with the seismic ambient noise such that evolution of seismic waves speed. This noise is almost all the signal continuously recorded by the seismic stations worldwide (Europe, China, USA, Japan), except earthquakes. It offers new observables for the seismologists, new types of virtual seismograms that are not only located at the place of earthquakes. For instance, one can obtain wave paths that probes the deepest part of the Earth.

Accordingly, this is one of the first project in the seismological community that studies systematically the continuous recordings, which represents a large amount of seismological data, of the order of several tens of terabytes. For instance, one year of the Japanese Network is about 20 TB or 3 months of the mobile network USArray represents 500 GB (it depends on the sampling of the recorders).

In addition, the calculation operations downstream may produce even more data than the observation data. To give an order of magnitude, more than 200 TB has been managed by the Whisper project at the same time. A classical processing produces 8 TB in 5 days. An other one 'read' 3 or 4 TB and 'produced' 1 TB in 6 hours. Many tests of the signal processing are done and computational data are deleted as and when required.

Nowadays, the earth sciences or more generally, sciences of the universe are widely engaged in data-intensive processing. This leads to design scientific workflow, towards data-intensive discovery and e-Science.

Reflected by the Whisper project, we have to organize the science objectives with the computer constraints. We have to take into account the duration of postdocs and PhD theses, as well as the availability of computer infrastructures and their ease of access. This leads to many questions about software development, including the genericity of computer code and the technical support. But it also influences in terms of choice of appropriate infrastructures.

Even if this project has his own resources, such a problem of data-intensive requires specific tools able to organize distributed data management and access to computational resources: a data grid environment.

The University of Grenoble offers, thanks to the High Performance Computing (HPC) centre *Ciment*, this kind of environment with the distributed file system *Irods* and the middleware *CiGri*.

It is thanks to the close collaboration between IT resources of Whisper and the infrastructures of the University that this project has been implemented as we show below.

---

\*FP7 ERC Advanced grant 227507, see [whisper.obs.ujf-grenoble.fr](http://whisper.obs.ujf-grenoble.fr)

### 3 Software for data-intensive processing (6-12§, 1-2 pages)

A part of the Whisper project is specifically dedicated to the IT codes. This includes to design a specification, an implementation and some optimisations of a sequence of a data-management and of a computations<sup>†</sup>. This project uses some own IT resources (servers, dedicated bay) but also uses common IT infrastructure of the university. We developed also adaptations for the IT infrastructures and we provide technical support for researchers.

Most of the IT codes are writing with the *Python* language and uses intensively the scientific libraries *Scipy* (fortran and C embeded) and *Obspy*<sup>‡</sup> (essentially the 'read' function). The latest one is dedicated to the seismological community.


The IT codes consists of several tools described schematically at the figure 2 and grouped into three parts. The first one concerns the signal processing, the second part permits the computation of the correlations and the last part consists of codes for the analysis of the correlations.

A first package provides a flexible way to process raw data, to specify a pipeline of pre-processing of signal. The user starts by specifying a directory, a set of seismic stations and a set of dates. Then the codes scan the directory and extracts all pieces of seismograms (also called traces) and rearranges them in a specific architecture of files in order to calculate the correlations to the next step. We use here intensively the function 'read' of the library *Obspy* which allows to open most of format files seismogram. The user also define his own sequence of processings. He can use the functions predefined but also the Python libraries he needs. Moreover he can add eventually his own processing.

The second package concerns correlations. Roughly speaking, a correlation is an operation with two seismograms that provides the coherence between the both signal (for a given time windows). Moreover, it represents, in

some favorables cases, a new virtual seismogram (it converges to the Green's function). Thus, the code computes all the correlations and provides an architecture of files that corresponds to all the couples of seismograms (for each date).

Figure 1: Step of the correlations



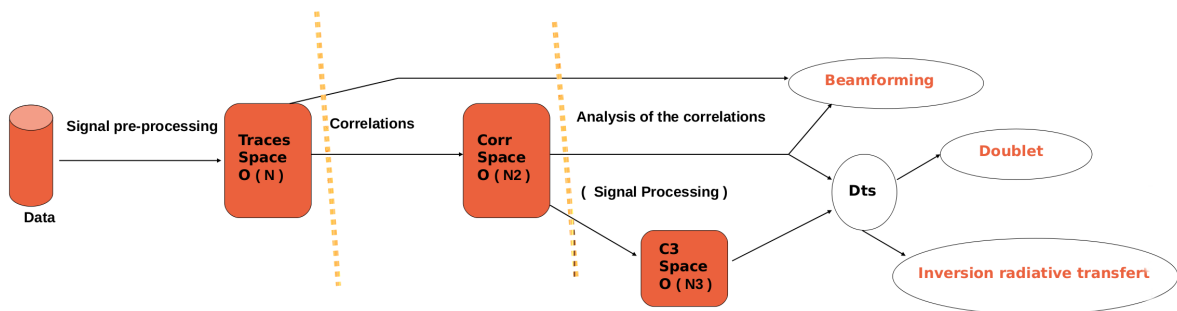
These quadratic space complexity can be critical and lot of effort was made in order to optimize the computation in two direction. First we improve the computation of the fast fourier transform by pre-calculating some "good" combinations of small primes numbers. With this method, we improve of forty percent the time computation in the favorable cases.

Nevertheless, the main optimization was made by testing the behaviour of the carbage collector of Python in order to follow the cache heuristics. More precisely, we do not use the gc module or the 'del' statement but we try to schedule and localize the line of code in order to find the good unfolding that uses the architecture optimally.

The last part of computer codes concerns the analysis of correlations (the virtual new seismograms) with methods such as beamforming, doublet or inversion. We also compute correlations of correlations C3 (also new seismograms). For example, we study the variations in velocity of seismic waves as we illustrate below.

These codes permit to process a dataset on a computer laptop. Nevertheless, to take advantage of IT infrastructure at the University of Grenoble, adjustments have been made for the grid computing as we shall see later.

Figure 2: Main sequences of processings of the Whisper Codes



<sup>†</sup> see [code-whisper.isterre.fr/html/](http://code-whisper.isterre.fr/html/) (part of the design)

<sup>‡</sup> see [www.obspy.org](http://www.obspy.org)

## 4 IT infrastructure for grid computing (6-12§, 1-2 pages)

### 4.1 CiGri and iRODS synergy

The data-intensive processing needs obviously an IT infrastructure in order to couple storage and computation. In our cases, most of the processing are embarrassingly parallel. The amount of data and the location of compute nodes available suggests using a distributed storage system and a grid manager.

The IT infrastructure used here is provided by the *Ciment* § platform. Ciment is the High Performance Computing (HPC) center of the Grenoble University. It offers a partial pooling of computing resources (10 computing clusters, 6600 cores and some GPUS) and many documentations for users. Moreover, the computational resources are integrated in a local grid of supercomputers. Associated with a distributed storage, it provides a local data grid environment.

The distributed storage accessible by all the computing nodes of all the clusters is managed by iRODS ¶. Nowadays it represents approximately 700 TB. The grid computing is managed by the *CiGri* || *middleware*, that is part of the OAR \*\* project (the Resource and Job Management System on which CiGri relies). CiGri and iRODS together build a complementary solution for embarrassingly parallel computations with large input/output distributed data sets.

Furthermore, with unitary parametric jobs that are reasonably short in time, CiGri can deal with the best-effort mode provided by OAR. In this mode, grid jobs are scheduled on free resources with a zero priority and may be killed at any time when the local demand of resources increases. This CIMENT organization (independent computing clusters glued together with a best-effort grid middleware and a distributed storage), in place for more than a decade, has proven to be very efficient, allowing near one hundred percent usage of computing resources thanks to small jobs being managed at the grid level.

Furthermore, as the results of the grid jobs are stored into the distributed storage with a unique namespace, iRODS also acts to the user as a centralized controller with a total observation and thus allows the user to monitor its calculation.


### 4.2 iRODS infrastructure

The Integrated Rule-Oriented Data System (iRODS) is a Distributed File System (DFS) offering a single namespace for files that are stored on different resources that may be on different locations. The administrator can set up

rules (microservices) to perform some automatic actions, for example the storage of a checksum or an automatic replication to the nearest resource (staging). The user can control himself replications and create user-defined meta-data. iRODS exposes a Command Line Interface (the i-commands), an API useable from several programming languages (C, python, PHP,...), a fuse interface, a web gui, and a webdav interface. The meta-catalog is an SQL database, which makes it very efficient for managing additional meta-data or making advanced queries (see [1] for an illustration of use). It is not "block-oriented", and thus relies on underlying Posix filesystems. Performance is not the main goal, but when the files are distributed on different resources, the only bottleneck is the meta-catalog (which is centralized).


The iRODS storage infrastructure of Ciment consists of a unique zone with the iCat server and a dozen of nodes as illustrated on figure 3. The nodes are grouped inside 3 different locations, called site A, site B and site C, having heterogeneous WAN connexions. Each site has its own 10Gbe local network switch. Those 3 sites are located into the 3 main datacenters where the CIMENT supercomputers live, so there are always close storage resources with the computing nodes. Note that site C has only a 1Gbe WAN connexion while A and B have 10Gbe WAN connexions.

Figure 3: iCAT and storage nodes



Capacity has now reached 700 TB and is constantly evolving and increases with investment in new projects, as iRODS offers a great scalability by simply adding new storage resources. Each node has currently 2 RAID arrays from 24 to 48 raw TB as illustrated at the figure 4

Figure 4: A storage node



CIMENT has set up a web interface where the user can easily check the status of the resources (figure 5).

Figure 5

§ see [ciment.ujf-grenoble.fr](http://ciment.ujf-grenoble.fr)


¶ see [irods.org](http://irods.org)

|| see [ciment.ujf-grenoble.fr/cigri/dokuwiki](http://ciment.ujf-grenoble.fr/cigri/dokuwiki)

\*\* see [oar.imag.fr](http://oar.imag.fr)

#### CIMENT STORAGE nodes states

Cigri (20 TB / resource)



### 4.3 CiGri infrastructure

The access to 6600 cores of the clusters of the Ciment platform is achieved through the middleware CiGri. CiGri launches embarrassingly parallel jobs on idle processors of every computing clusters and then optimizes the resources usage which are used for parallel jobs otherwise.

Each cluster of the University of Grenoble uses the resource manager OAR. CiGri acts, among other things, as a metascheduler of OAR. It retrieves the clusters states trough OAR and submits the jobs on free resources without exhausting the local scheduling queues.

While it may work in normal mode, CiGri is mostly used in best-effort mode and thus provides an automatic resubmission mechanism. CiGri offers a customizable notification system with a smart events management. With those mechanisms, the user can submit a big amount of small jobs, called a campaign, and forget about it until all the jobs of the campaign are terminated or CiGri notifies a serious problem.

Roughly speaking, in order to run a campaign, the user describes through a file (in the JSON format) the parameters of the campaign such as the accepted cluters, the needed resources, the maximum duration, the location of the codes, a prologue or epilogue script,... Codes and input data are retrieved from iRODS using i-commands into the prologue scripts and the jobs scripts (or using the iRODS API if the jobs are written into a supported language). So, there's no direct connexion between CiGri and iRODS, but the usage is totally complementary through the jobs scripts. Moreover, the user defines also a file where each line represents a value of the parameter for the user's code. Thus, the number of line of these parameter file corresponds to the number of jobs of the campaign.

CiGri is now at the version 3, which represents a major evolution in terms of modeling and technology (Rest API, Ruby)

## 5 Results (28-31§, 6-7 pages)

### References

- [1] Gen-Tao Chiang, Peter Clapham, Guoying Qi, Kevin Sale and Guy Coates, Implementing a genomic

data management system using iRODS in the Wellcome Trust Sanger Institute. *BMC Bioinformatics*, 12(1):361+, September 2011.